Butterfly Effects of SGD Noise:
Error Amplification in Behavior Cloning and Autoregression [1]

Adam Block♦, Dylan Foster♠, Akshay Krishnamurthy♠, Max Simchowitz♦, Cyril Zhang♠ [2]
♦MIT ♠Microsoft Research NYC

Behavior cloning & feedback loops [3]

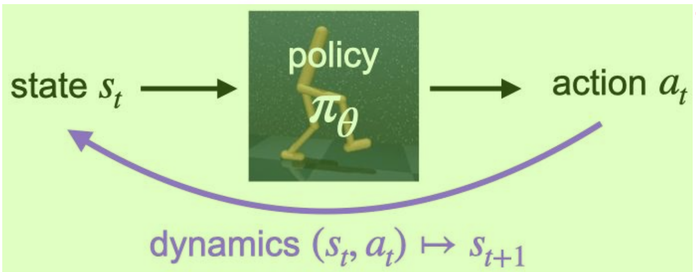

[6]

BC: fit θ to a collection of demonstrations from expert policy $\pi^\star$, minimizing 1-step loss [7]

(e.g.) $\ell_{\mathrm{BC}}(\theta) = \mathop{\mathbb{E}}_{\text{expert's } (s_t, a_t)} \left[\|\pi_\theta(s_t) - \pi^\star(s_t)\|^2 \right]$ [8]

Usual view: BC reduces offline imitation learning to supervised learning [9]

$\ell_{\mathrm{BC}}(\theta) \le \varepsilon$ (train on 1-step loss) → $J_H(\pi_\theta) \ge J_H(\pi^\star) - C \cdot \varepsilon$ (evaluate multi-step rollout rewards) [10]

Q1: *When BC works, how big is C in practice?* [11]

Q2: *When BC fails, must we improve the data? Or can we improve upon standard training?* [12]

Special case: pretraining vs. generation in LLMs [13]

- s_t : sequence of t tokens
- a_t : append a single token
- π_θ : autoregressive language model [14]

Gradient Variance Amplification [4]

A1: *Long-horizon rollouts can be extremely sensitive to small SGD fluctuations.* [15]

GVA: an empirical pathology in BC with NNs [16]

- *small changes in θ due to minibatch noise* [17]
- *...causing small changes in 1-step $\ell_{\mathrm{BC}}(\theta)$...* [18]
- *...and **chaotic** long-rollout rewards $J_H(\pi_\theta)$* [19]

Evidence: *(across locomotion tasks & archs)* [20]

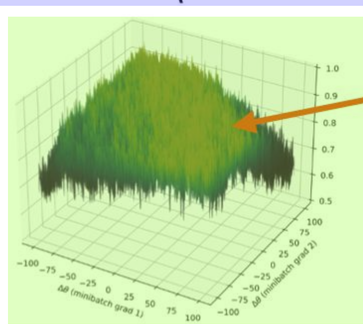

[22]

In stochastic gradient directions: [21]
- small changes in ℓ_{BC}
- large changes in Walker2D reward

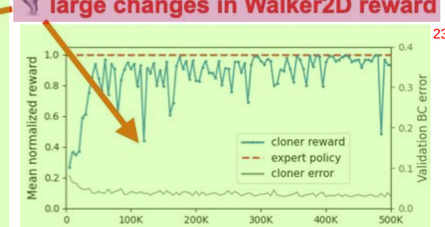

[23]

fractal reward landscape **oscillations during training** [24]

(in the paper: accompanying quantitative studies)

Theory: vignettes in toy models [25]

- marginally stable linear dynamics, linear $\pi, \pi^\star$ [26]
- "CliffLQR": halfspace reward → GVA
- SDE limit: EMA & LR schedules both work

Algorithmic mitigations for GVA [5]

A2: *Variance reduction techniques mitigate the oscillations; EMA is highly effective.* [27]

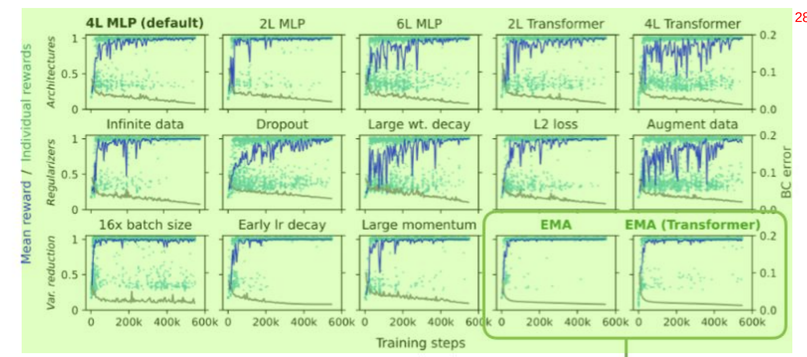

[28]

Exponential Moving Average [29]

$\theta_{t+1} \leftarrow \theta_t - \eta_t \widetilde{\nabla}_\theta \ell_{\mathrm{BC}}(\theta)$ (or Adam, etc.)

$\bar{\theta}_{t+1} \leftarrow (1-\beta_t)\bar{\theta}_t + \beta_t \cdot \theta_t$, use $\bar{\theta}_t$ for inference [30]

Bonus: GVA & EMA in language generation [31]

[33]

[32] [34]

EMA improves & stabilizes LM iterates before learning rate cooldown [35]